

setitimer(2) - Linux man page

Name

getitimer, setitimer - get or set value of an interval timer

Synopsis

```
#include <sys/time.h>
```

```
int getitimer(int which, struct itimerval *curr_value);
int setitimer(int which, const struct itimerval *new_value,
              struct itimerval *old_value);
```

Description

The system provides each process with three interval timers, each decrementing in a distinct time domain. When any timer expires, a signal is sent to the process, and the timer (potentially) restarts.

ITIMER_REAL

decrements in real time, and delivers **SIGALRM** upon expiration.

ITIMER_VIRTUAL

decrements only when the process is executing, and delivers **SIGVTALRM** upon expiration.

ITIMER_PROF

decrements both when the process executes and when the system is executing on behalf of the process. Coupled with **ITIMER_VIRTUAL**, this timer is usually used to profile the time spent by the application in user and kernel space. **SIGPROF** is delivered upon expiration.

Timer values are defined by the following structures:

```
struct itimerval {
    struct timeval it_interval; /* next value */
    struct timeval it_value;    /* current value */
};

struct timeval {
    time_t      tv_sec;          /* seconds */
    suseconds_t tv_usec;        /* microseconds */
};
```

The function **getitimer()** fills the structure pointed to by *curr_value* with the current setting for the timer specified by *which* (one of **ITIMER_REAL**, **ITIMER_VIRTUAL**, or **ITIMER_PROF**). The element *it_value* is set to the amount of time remaining on the timer, or zero if the timer is disabled. Similarly, *it_interval* is set to the reset value.

The function **setitimer()** sets the specified timer to the value in *new_value*. If *old_value* is non-NULL, the old value of the timer is stored there.

Timers decrement from *it_value* to zero, generate a signal, and reset to *it_interval*. A timer which is set to zero (*it_value* is zero or the timer expires and *it_interval* is zero) stops.

Both *tv_sec* and *tv_usec* are significant in determining the duration of a timer.

Timers will never expire before the requested time, but may expire some (short) time afterward, which depends on the system timer resolution and on the system load; see [time\(7\)](#). (But see **BUGS** below.) Upon expiration, a signal will be generated and the timer reset. If the timer expires while the process is active (always true for **ITIMER_VIRTUAL**) the signal will be delivered immediately when generated. Otherwise the delivery will be offset by a small time dependent on the system loading.

Return Value

On success, zero is returned. On error, -1 is returned, and *errno* is set appropriately.

Errors

EFAULT

new_value, *old_value*, or *curr_value* is not valid a pointer.

EINVAL

which is not one of **ITIMER_REAL**, **ITIMER_VIRTUAL**, or **ITIMER_PROF**; or (since Linux 2.6.22) one of the *tv_usec* fields in the structure pointed to by *new_value* contains a value outside the range 0 to 999999.

Conforming To

POSIX.1-2001, SVr4, 4.4BSD (this call first appeared in 4.2BSD). POSIX.1-2008 marks **getitimer()** and **setitimer()** obsolete, recommending the use of the POSIX timers API (**timer_gettime(2)**, **timer_settime(2)**, etc.) instead.

Notes

A child created via **fork(2)** does not inherit its parent's interval timers. Interval timers are preserved across an **execve(2)**.

POSIX.1 leaves the interaction between **setitimer()** and the three interfaces **alarm(2)**, **sleep(3)**, and **usleep(3)** unspecified.

The standards are silent on the meaning of the call:

```
setitimer(which, NULL, &old_value);
```

Many systems (Solaris, the BSDs, and perhaps others) treat this as equivalent to:

```
getitimer(which, &old_value);
```

In Linux, this is treated as being equivalent to a call in which the *new_value* fields are zero; that is, the timer is disabled. *Don't use this Linux misfeature*: it is nonportable and unnecessary.

Bugs

The generation and delivery of a signal are distinct, and only one instance of each of the signals listed above may be pending for a process. Under very heavy loading, an **ITIMER_REAL** timer may expire before the signal from a previous expiration has been delivered. The second signal in such an event will be lost.

On Linux kernels before 2.6.16, timer values are represented in jiffies. If a request is made set a timer with a value whose jiffies representation exceeds **MAX_SEC_IN_JIFFIES** (defined in *include/linux/jiffies.h*), then the timer is silently truncated to this ceiling value. On Linux/i386 (where, since Linux 2.6.13, the default jiffy is 0.004 seconds), this means that the ceiling value for a timer is approximately 99.42 days. Since Linux 2.6.16, the kernel uses a different internal representation for times, and this ceiling is removed.

On certain systems (including i386), Linux kernels before version 2.6.12 have a bug which will produce premature timer expirations of up to one jiffy under some circumstances. This bug is fixed in kernel 2.6.12.

POSIX.1-2001 says that **setitimer()** should fail if a *tv_usec* value is specified that is outside of the range 0 to 999999. However, in kernels up to and including 2.6.21, Linux does not give an error, but instead silently adjusts the corresponding seconds value for the timer. From kernel 2.6.22 onward, this nonconformance has been repaired: an improper *tv_usec* value results in an **EINVAL** error.

See Also

gettimeofday(2), **sigaction(2)**, **signal(2)**, **timer_create(2)**, **timerfd_create(2)**, **time(7)**

Referenced By

__pmafisempty(3), **perlfunc(1)**, **profil(3)**, **pthread(7)**, **signal(7)**, **snmp_alarm(3)**, **socket(7)**, **ualarm(3)**